

SPARK TP2 Soufiane MOUTEI

October 19, 2018

1 Let's start with initializing Spark:

```
In [1]: from pyspark import SparkContext
        from pyspark.sql import SQLContext
        from pyspark.sql.types import *
        from pyspark.sql.functions import udf, first

        sc = SparkContext()
        sqlContext = SQLContext(sc)
```

```
In [2]: sc
```

```
Out[2]: <SparkContext master=local[*] appName=pyspark-shell>
```

2 Average monthly income of the shop in France:

```
In [3]: # A function to get: store, month, income; it will be used for mapping
        def f(x):
            key = x[0]
            line = x[1].split()
            return key, line[0], int(line[1])
```

```
In [4]: # Let's start with preparing the files
        whole_file = (
            # read the text files of the directory "input1"
            sc.wholeTextFiles("input1/")

            # Only the key is changed here: we split on "/" to get the textfile name
            # and then delete the file extension
            .map(lambda x: (x[0].split("/)[-1][: -4], x[1]))

            # Split the values to get each line on the file linked to the city name
            .flatMapValues(lambda x: x.split("\r\n"))

            # Map the function f to get: city, month, income.
            .map(f)
```

```

)

# The schema that will be used for 'whole_file'
mySchema = StructType([
    StructField("Store", StringType(), True),
    StructField("Month", StringType(), True),
    StructField("Income", IntegerType(), True)
])

# Create a DataFrame using 'whole_file' and the schema 'mySchema'
df = sqlContext.createDataFrame(whole_file, schema=mySchema)

# Add the city name as a column to 'df'
getCity = udf(lambda t: t.split("_")[0])
df = df.withColumn("City", getCity(df["Store"]))

```

In [5]: # Answer to the question

```

monthly_income_france = (
    df

    # To select as columns: Month, Income
    .select("Month", "Income")

    # Group by 'Month' and take the average of 'Income' for each month
    .groupBy("Month")
    .avg()

    # Rename the column representing the average of 'Income' for each month
    .withColumnRenamed('avg(Income)', 'Average monthly income')
)
monthly_income_france.show()

```

```

+-----+-----+
|Month|Average monthly income|
+-----+-----+
| APR|      20.23076923076923|
| OCT|      26.53846153846154|
| NOV|      24.53846153846154|
| FEB|      19.153846153846153|
| SEP|      25.53846153846154|
| JAN|      20.76923076923077|
| AUG|      23.076923076923077|
| MAR|      17.53846153846154|
| DEC|               29.0|
| JUN|      27.846153846153847|
| JUL|      21.692307692307693|
| MAY|      22.46153846153846|
+-----+-----+

```

3 Average monthly income of the shop in each city:

```
In [6]: monthly_income_per_city = (  
    df  
  
    # To select as columns: City, Month, Income  
    .select("City", "Month", "Income")  
  
    # Group by 'City' and 'Month' and take the average of 'Income' for each group  
    .groupBy("City", "Month")  
    .avg()  
  
    # Rename the column representing the average of 'Income'  
    .withColumnRenamed('avg(Income)', 'Average monthly income')  
  
    # Sort by the city to get the values of each city together  
    .orderBy('City', ascending=True)  
)  
monthly_income_per_city.show(200)
```

City	Month	Average monthly income
anger	NOV	14.0
anger	AUG	15.0
anger	JUL	19.0
anger	JUN	15.0
anger	JAN	13.0
anger	DEC	16.0
anger	OCT	8.0
anger	FEB	12.0
anger	SEP	13.0
anger	MAR	14.0
anger	APR	15.0
anger	MAY	12.0
lyon	MAR	14.0
lyon	AUG	25.0
lyon	APR	15.0
lyon	OCT	11.0
lyon	JUN	15.0
lyon	MAY	12.0
lyon	JUL	19.0
lyon	NOV	22.0
lyon	SEP	13.0

	lyon	FEB	12.0
	lyon	DEC	22.0
	lyon	JAN	13.0
	marseilles	SEP	23.0
	marseilles	DEC	26.0
	marseilles	NOV	24.0
	marseilles	APR	22.0
	marseilles	JUN	25.0
	marseilles	OCT	28.0
	marseilles	JAN	16.0
	marseilles	AUG	22.0
	marseilles	MAY	18.5
	marseilles	JUL	21.0
	marseilles	MAR	16.0
	marseilles	FEB	16.0
	nantes	APR	12.0
	nantes	AUG	11.0
	nantes	MAY	21.0
	nantes	FEB	15.0
	nantes	NOV	14.0
	nantes	OCT	14.0
	nantes	JUN	28.0
	nantes	DEC	24.0
	nantes	JAN	16.0
	nantes	SEP	13.0
	nantes	MAR	20.0
	nantes	JUL	19.0
	nice	JUL	19.0
	nice	MAY	11.0
	nice	MAR	20.0
	nice	JAN	16.0
	nice	FEB	15.0
	nice	OCT	18.0
	nice	JUN	18.0
	nice	AUG	11.0
	nice	DEC	29.0
	nice	NOV	14.0
	nice	APR	9.0
	nice	SEP	23.0
	orlean	OCT	8.0
	orlean	MAR	14.0
	orlean	APR	15.0
	orlean	NOV	24.0
	orlean	FEB	12.0
	orlean	SEP	13.0
	orlean	JUL	19.0
	orlean	JUN	15.0
	orlean	AUG	25.0

	orlean	MAY	12.0
	orlean	JAN	13.0
	orlean	DEC	26.0
	paris	JUL	33.66666666666664
	paris	OCT	56.66666666666664
	paris	JAN	38.333333333333336
	paris	FEB	33.0
	paris	SEP	48.0
	paris	MAR	26.333333333333332
	paris	MAY	50.0
	paris	NOV	48.66666666666664
	paris	DEC	52.66666666666664
	paris	JUN	55.0
	paris	AUG	41.66666666666664
	paris	APR	38.66666666666664
	rennes	AUG	11.0
	rennes	APR	9.0
	rennes	JUN	13.0
	rennes	FEB	18.0
	rennes	SEP	23.0
	rennes	DEC	20.0
	rennes	JUL	14.0
	rennes	OCT	18.0
	rennes	MAR	10.0
	rennes	JAN	19.0
	rennes	NOV	14.0
	rennes	MAY	11.0
	toulouse	APR	11.0
	toulouse	SEP	23.0
	toulouse	DEC	19.0
	toulouse	JAN	12.0
	toulouse	MAY	11.0
	toulouse	NOV	12.0
	toulouse	MAR	14.0
	toulouse	AUG	11.0
	toulouse	FEB	13.0
	toulouse	JUN	18.0
	toulouse	OCT	14.0
	toulouse	JUL	19.0
	troyes	JUN	25.0
	troyes	JUL	11.0
	troyes	DEC	11.0
	troyes	AUG	22.0
	troyes	FEB	21.0
	troyes	APR	17.0
	troyes	SEP	21.0
	troyes	MAR	11.0
	troyes	NOV	11.0

	troyes	OCT	28.0
	troyes	JAN	21.0
	troyes	MAY	15.0
+-----+			

4 Total revenue per city per year:

```
In [7]: yearly_income_city = (
        df

        # To select as columns: City, Income
        .select("City", "Income")

        # Group by 'City' and take the sum of 'Income' for each city
        .groupBy("City")
        .sum()

        # Rename the column representing the sum of 'Income'
        .withColumnRenamed('sum(Income)', 'Total revenue per year')
    )
    yearly_income_city.show(20)
```

+-----+	
	City Total revenue per year
+-----+	
	nantes 207
	troyes 214
	paris 1568
	lyon 193
	anger 166
	marseilles 515
	nice 203
	orlean 196
	rennes 180
	toulouse 177
+-----+	

5 Total revenue per store per year:

```
In [8]: yearly_income_store = (
        df

        # To select as columns: Store, Income
```

```

.select("Store", "Income")

# Group by 'Store' and take the sum of 'Income' for each store
.groupBy("Store")
.sum()

# Rename the column representing the sum of 'Income'
.withColumnRenamed('sum(Income)', 'Total revenue per year')
)
yearly_income_store.show(20)

```

Store	Total revenue per year
nantes	207
troyes	214
lyon	193
marseilles_1	284
paris_2	642
anger	166
paris_3	330
marseilles_2	231
nice	203
orlean	196
rennes	180
paris_1	596
toulouse	177

6 The store that achieves the best performance in each month:

```

In [9]: # Method 1: Using sorting
        best_performance_store = (

        # To select as columns: Month, Store, Income
        df.select("Month", "Store", "Income")

        # Sort by (month, -income) to get the values sorted by the month,
        # and in case we have the same month, we sort in a descending order the income
        .orderBy('Month', - df['Income'], ascending=True)

        # group by the month to get the best store for each month (that is the first store s
        .groupBy("Month")
        .agg(first("Store").alias("Best Store"))
        )
        best_performance_store.show(20)

```

```

+-----+-----+
|Month|Best Store|
+-----+-----+
|  APR|   paris_1|
|  OCT|   paris_1|
|  NOV|   paris_2|
|  FEB|   paris_2|
|  SEP|   paris_2|
|  JAN|   paris_1|
|  AUG|   paris_2|
|  MAR|   paris_2|
|  DEC|   paris_1|
|  JUN|   paris_2|
|  JUL|   paris_1|
|  MAY|   paris_2|
+-----+-----+

```

```
In [10]: # Method 2: Without sorting, using the maximum function
```

```
best_performance_store_2 = (
    df
```

```
    # To select as columns: Month, Income
    .select("Month", "Income")
```

```
    # group by the month to get the maximum income for each month
    .groupby("Month")
    .max("Income")
```

```
    # Join the initial data (selecting: Month, Store, Income) on the month
    .join(df.select("Month", "Store", "Income"), "Month", how='outer')
)
```

```
best_performance_store_2 = (
    best_performance_store_2
```

```
    # Select only where the income is equal to the maximum, which is the best value
    .filter(best_performance_store_2["Income"] == best_performance_store_2["max(Income)"]
```

```
    # Drop the maximum column and the income column because we need only the month and
    .drop("max(Income)", "Income")
)
```

```
best_performance_store_2.show(20)
```

```

+-----+-----+
|Month|  Store|
+-----+-----+

```


	APR		paris_1	
	MAY		paris_2	
	SEP		paris_2	
	JUN		paris_2	
	JAN		paris_1	
	OCT		paris_1	
	NOV		paris_2	
	FEB		paris_2	
	MAR		paris_2	
	AUG		paris_2	
	JUL		paris_1	
	DEC		paris_1	
+	-----	+	-----	+